LaTeX-Native Paper Figures

Learning a Symbolic Nonverbal Vocabulary from Dyadic Pose and Gaze Interactions

All diagrams, plots, labels, and values in this document are rendered directly by LaTeX using TikZ and PGFPlots. No raster image is embedded.

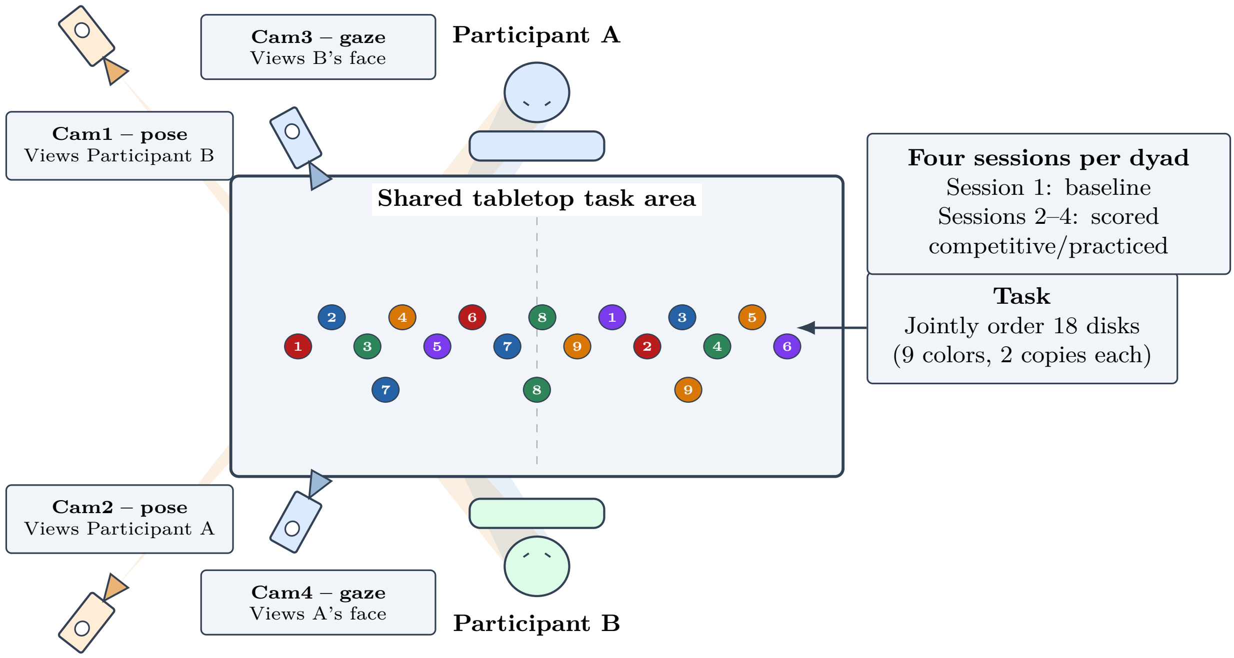

Figure 1: Experimental setup and task structure. Each camera icon is oriented toward the opposing participant, and its field-of-view cone terminates at the intended target. Pose cameras target the opposing upper body and face; gaze cameras target the opposing face.

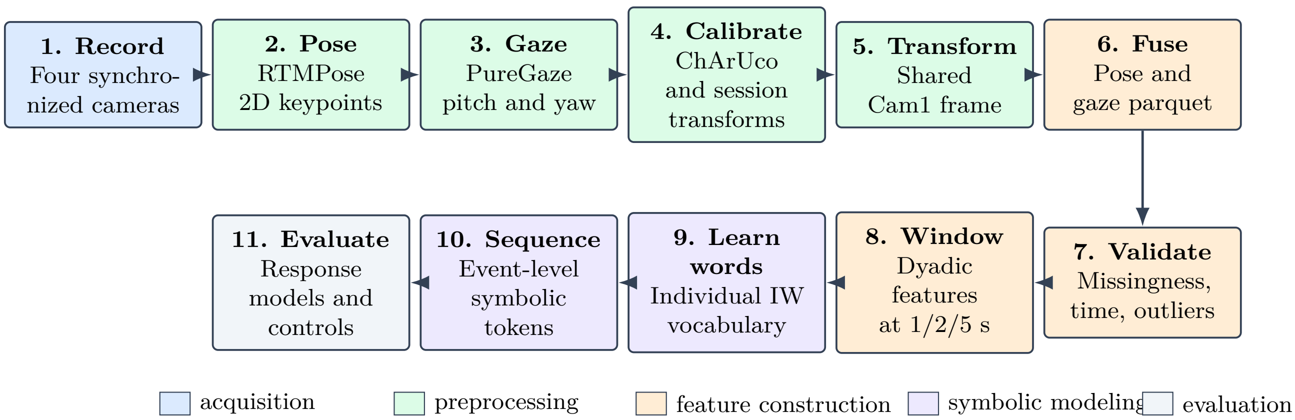

Figure 2: End-to-end analysis pipeline. The pipeline is arranged in two rows to keep every processing step and arrow readable at paper width.

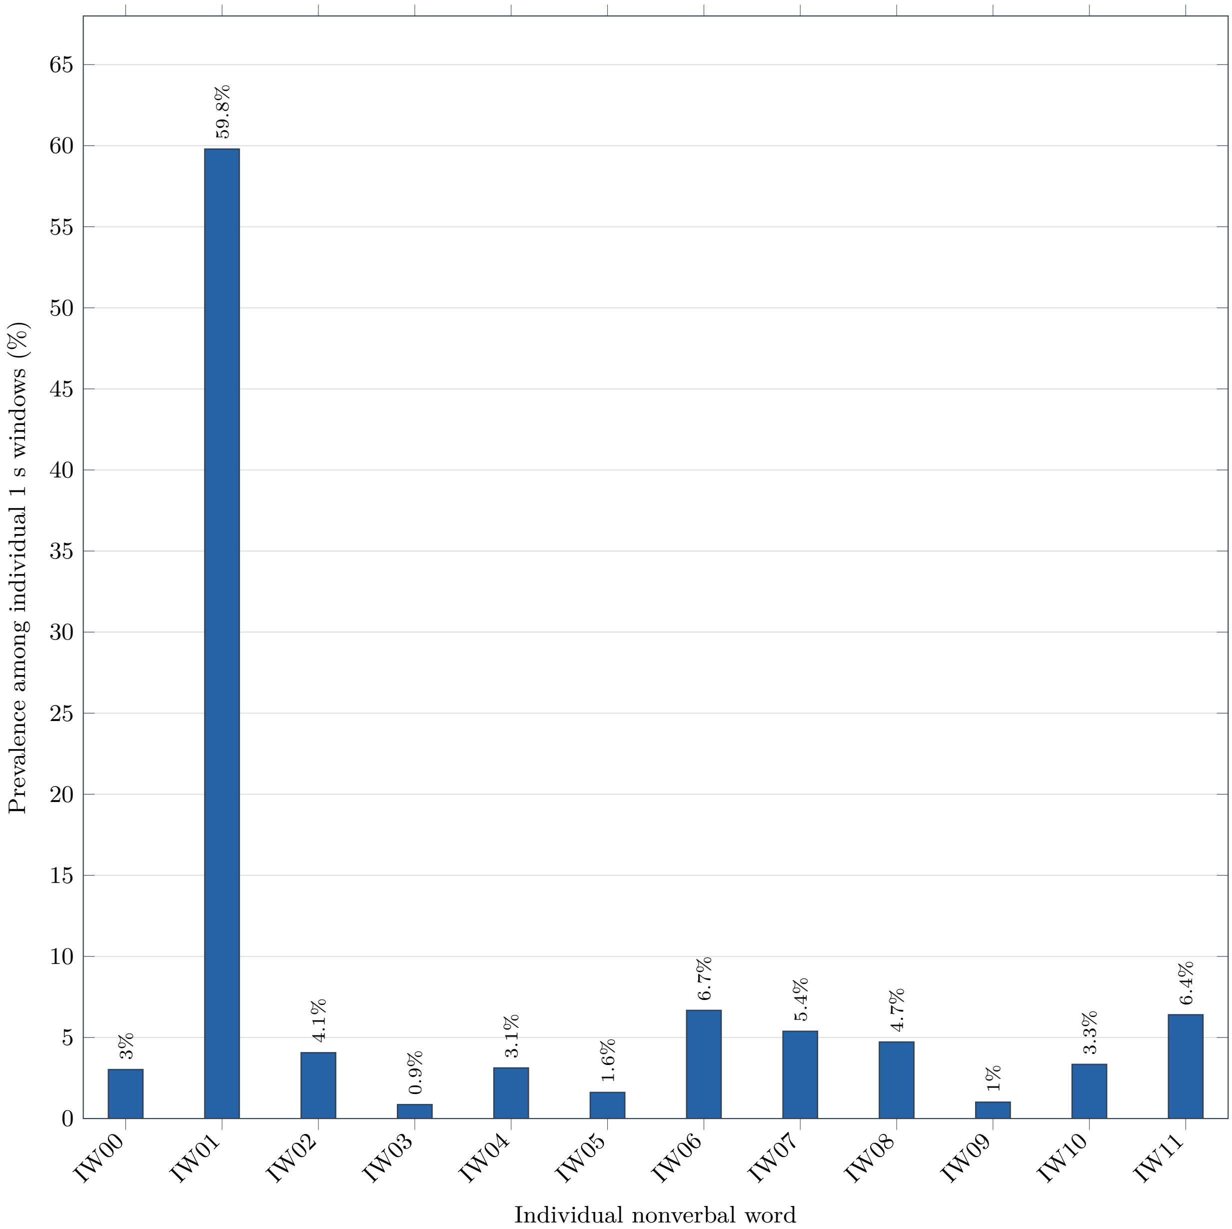

Figure 3: Prevalence of the learned individual nonverbal vocabulary. All bars use one color so magnitude, rather than category color, carries the visual comparison.

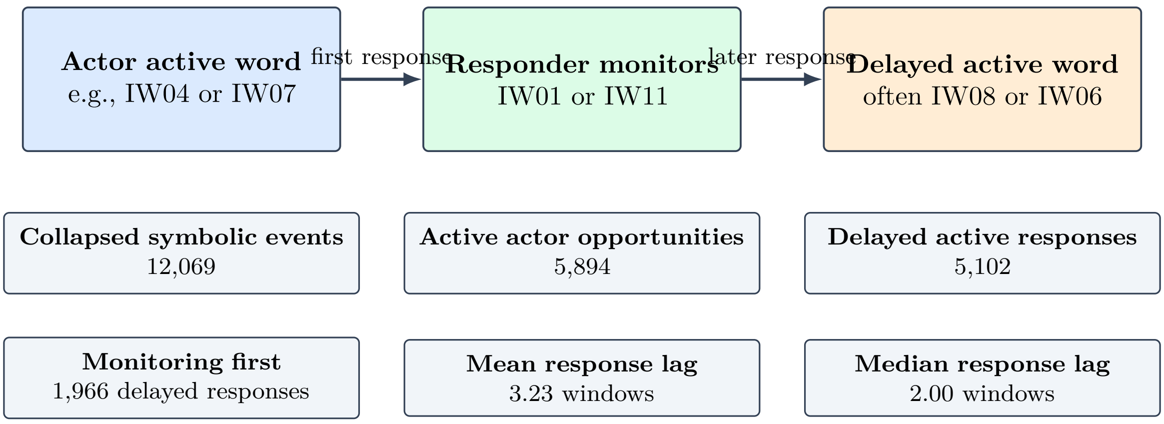

Figure 4: Event-level symbolic analysis revealed a recurring action–monitoring–action structure.

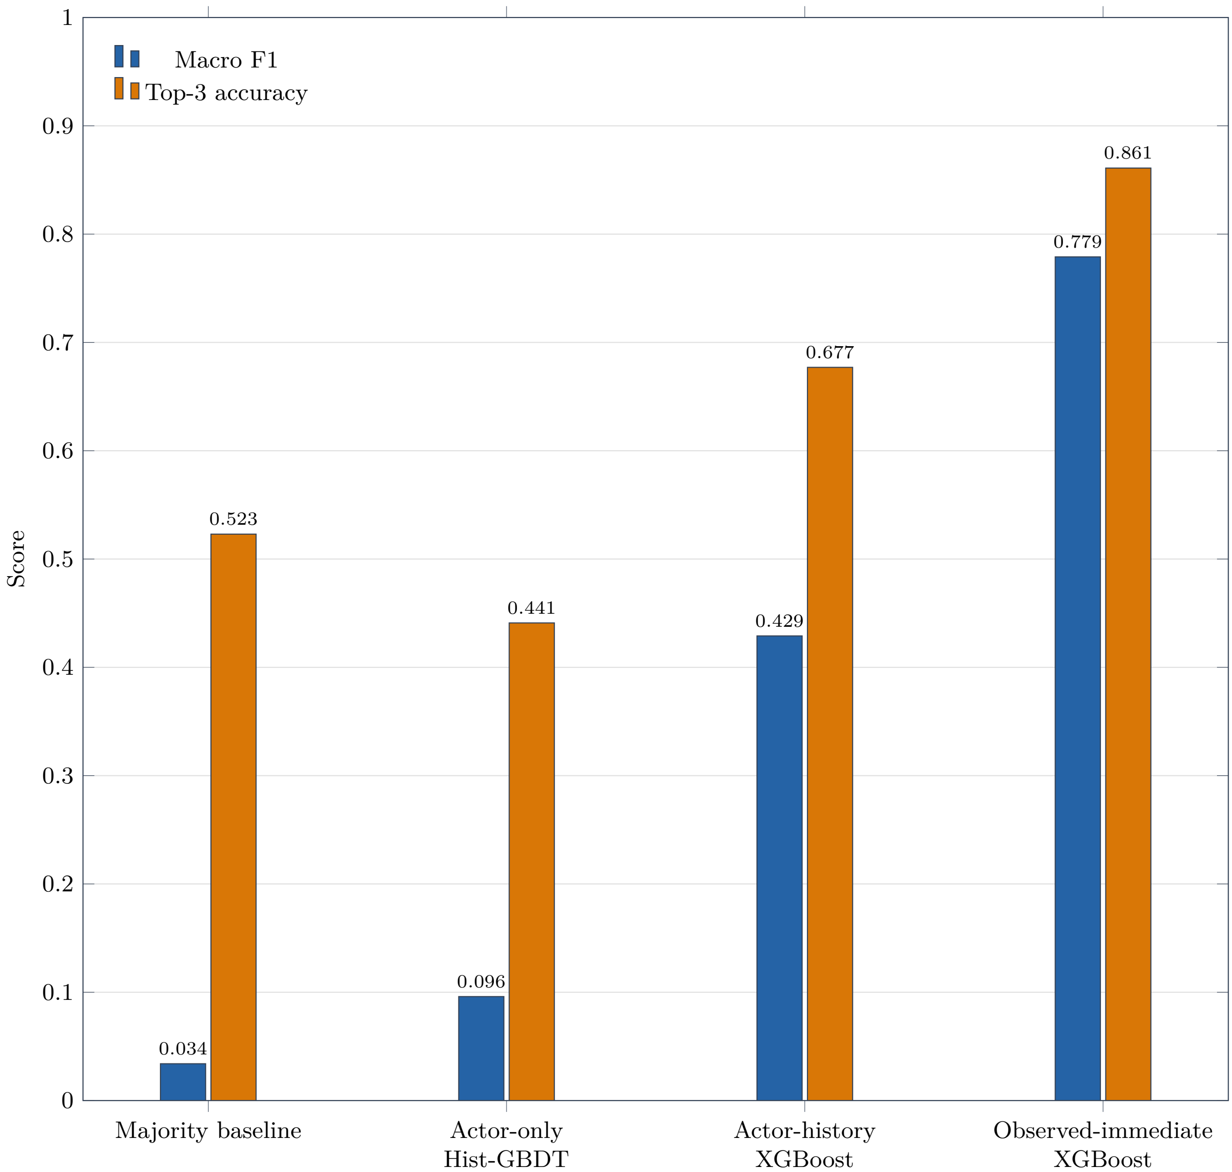

Figure 5: Event-level response modeling under leave-one-dyad-out evaluation. Observed-immediate models are suitable for offline interpretation because they observe the responder's first post-actor state.

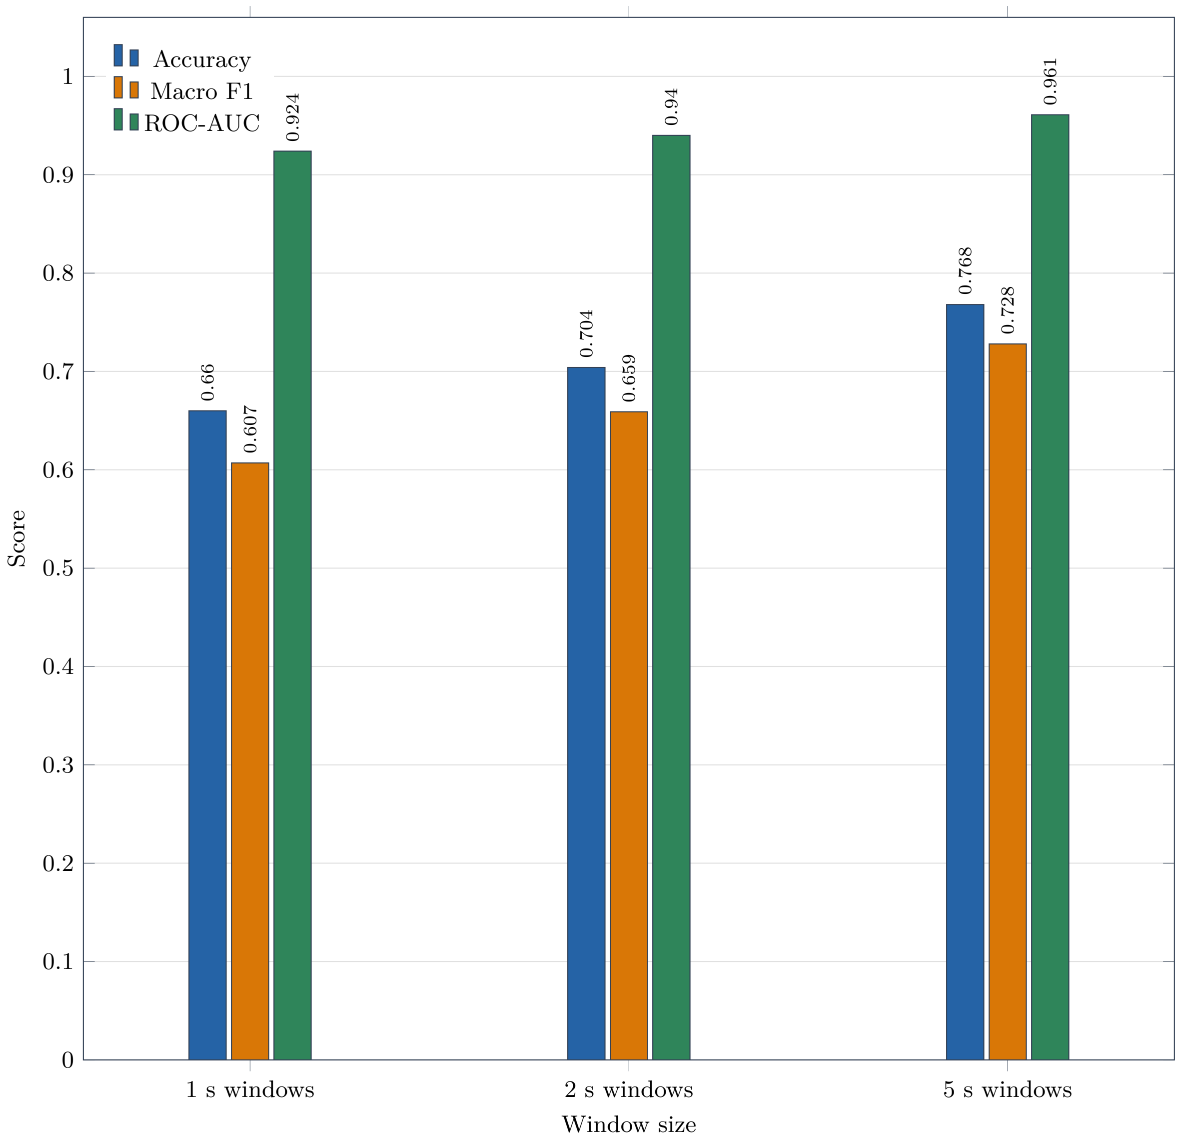

Figure 6: Pair-identity control experiment. High pair-recognition performance shows that dyads have recognizable signatures and supports leave-one-dyad-out evaluation.

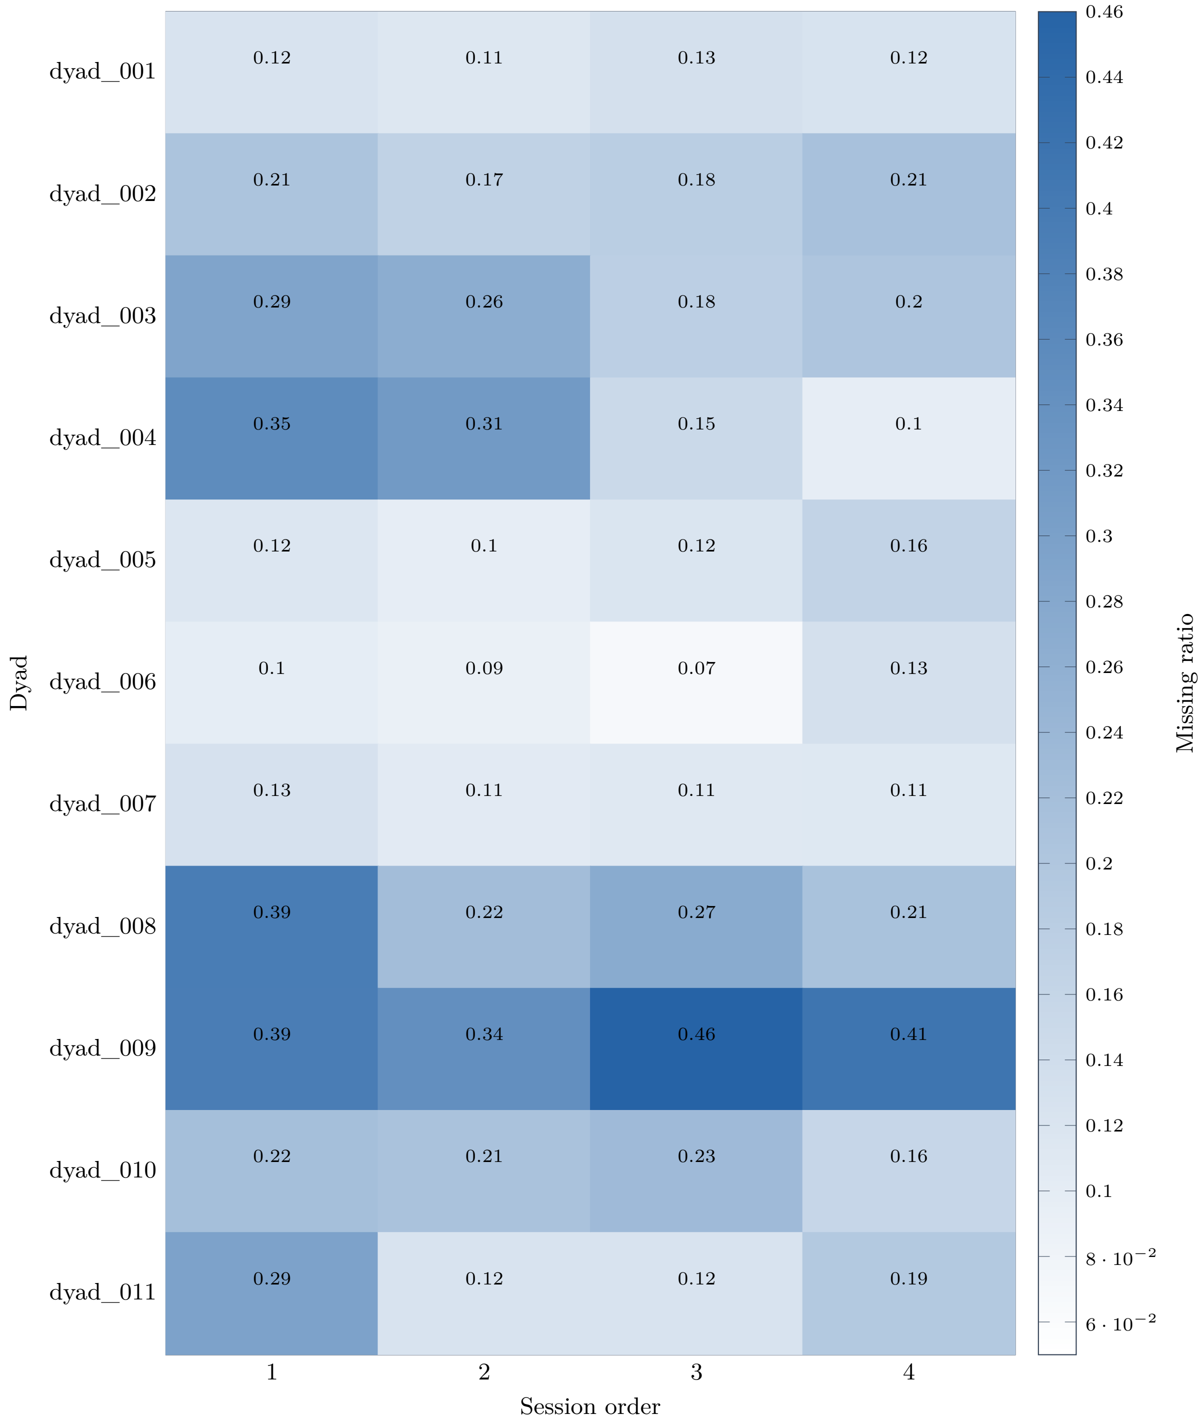

Figure A1: Session-level missingness aggregated by dyad and session order.

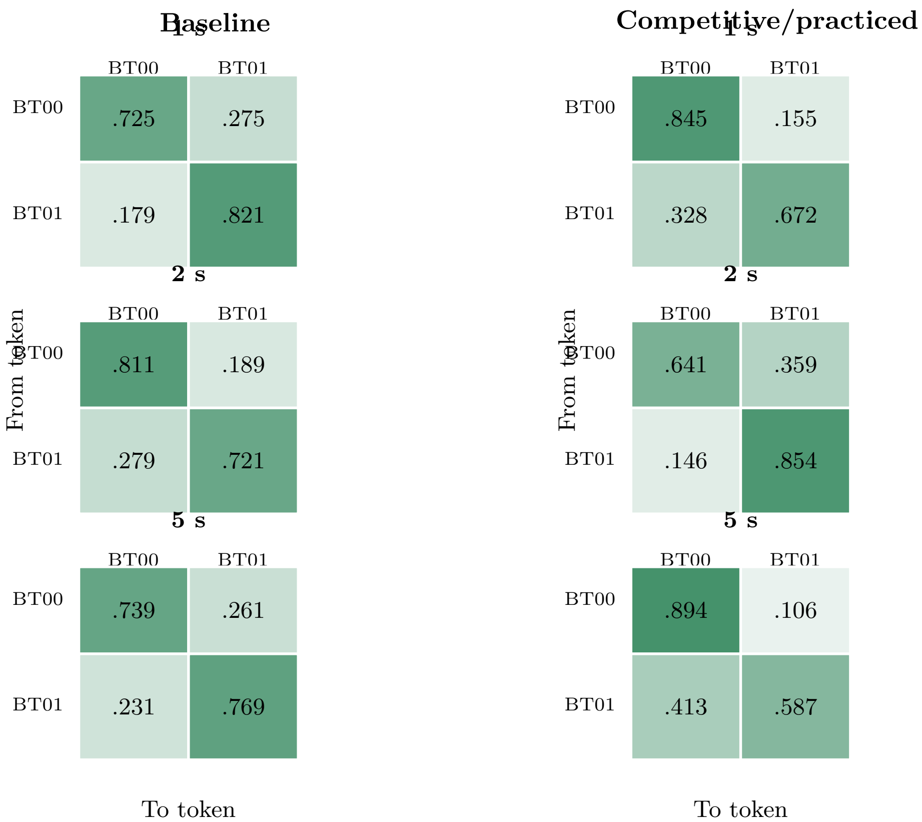

Figure A2: Baseline and competitive/practiced token transition matrices. Jensen–Shannon divergences are 0.0774 at 1 s, 0.0903 at 2 s, and 0.1081 at 5 s.

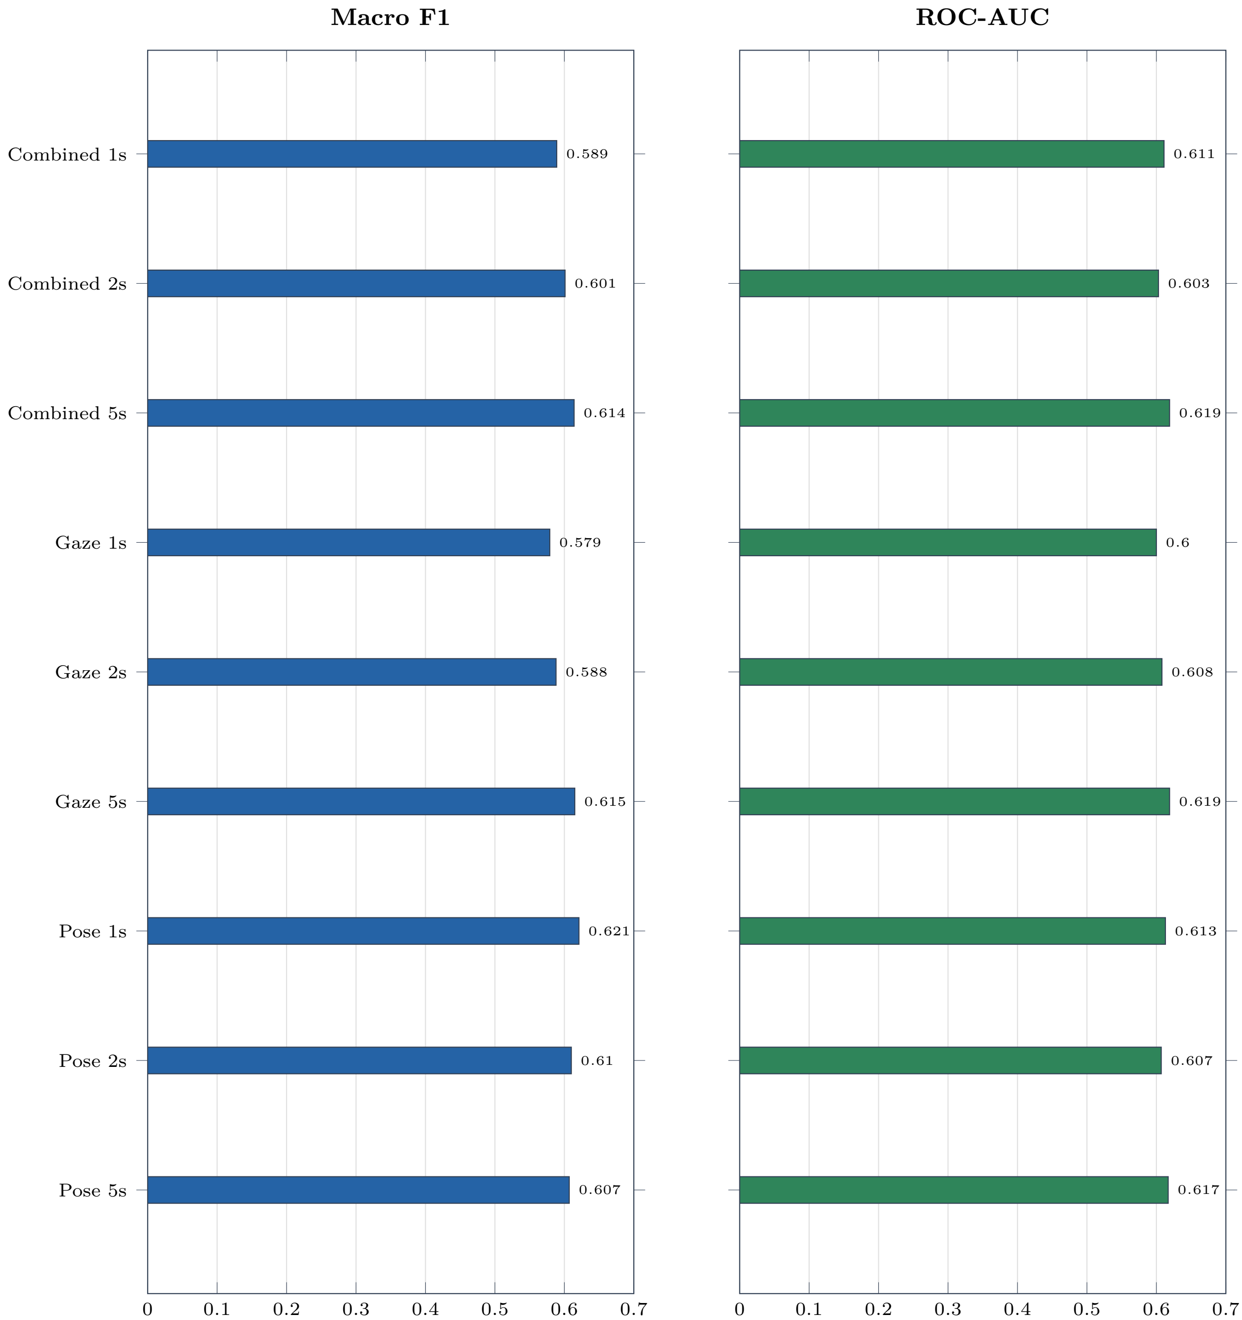

Figure A3: Best non-dummy baseline-vs-competitive/practiced random-forest classifier per window size and feature set.